

Tabular Database Systems

②

Tabular Data in CSV Files

Summer 2025

Torsten Grust
Universität Tübingen, Germany


1 | Most Data Lives Outside DBMSs

Reality check: The vast majority of data out there does *not* live in files and formats directly accessible for DBMSs:

- Applications read/write their specific data formats which were probably devised before DBMS X became available.
- Users process their data using a variety of systems. Why lock-in to the proprietary database format of DBMS X (or Y or 🐧)?
 - Rather build on a “lowest common denominator” format that allows **data interchange between systems**.
- Simplicity wins! DBMS-internal data formats are intricate, heavily optimized for access speed and space efficiency.
 - Rather use formats that admit **simple read/write routines**.
 - Let **humans read/write the data format**, ideally using no tools but their favorite **plain text** or spreadsheet editor.

2 : Tabular Data in Comma-Separated Value (CSV) Files ①

CSV files (*.csv) provide one such lowest common denominator DBMS-external dataformat. CSV encodes tabular data in **line-structured plain text** files:



1. Newline (NL) separates lines \equiv , a **line** encodes one table **row**.
2. Delimiters (often comma $,$) split lines into **columns**.
 - We expect each line to hold the same number of columns, leading to a rectangular grid of cells $\begin{smallmatrix} \blacksquare & \blacksquare & \blacksquare \\ \blacksquare & \blacksquare & \blacksquare \end{smallmatrix}$.

Tabular Data in Comma-Separated Value (CSV) Files ②

CSV file `007-vehicles.csv` to define the column names and contents of table `vehicles` (see Chapter ①):

```
vehicle, kind, seats, wheels?
```

```
🚗, car, 5, true
```

```
🚙, SUV, 3, true
```

```
🚌, bus, 42, true
```

```
🚒, bus, 7, true
```

```
🚲, bike, 1, true
```

```
🚓, tank, , false
```

```
🚗, cabrio, 2, true
```

📄 #007

- **Notes:**

- Column names held in first CSV line “by convention.”
 - Otherwise, **header and data rows are identical.**
- Columns are **untyped** (`wheels?` “looks like” type `boolean`).
- **NULL** represented by adjacent delimiters(`,` `,`).

Reading CSV Files Like Database Tables

- In DuckDB, read a CSV file much like a database-internal table:

```
D FROM '007-vehicles.csv';
```

vehicles	kind	seats	wheels?

- This is a shorthand for an explicit call to built-in function `read_csv()`:

```
D FROM read_csv('007-vehicles.csv', configuration parameters);
```











vehicles	kind	seats	wheels?

- Non-default *configuration parameters* may be required to disambiguate CSV file contents regarding presence of headers, choice of delimiters, column types, ...)

3 | CSV Enables Data Exchange

A variety of systems export data in CSV format. Paired with the DBMS's CSV import, this enables **CSV-based data exchange**.

- Example: Spreadsheets. On Google Sheets, use command **File ► Download ► Comma-Separated Values (.csv)**:

	A	B	C	D
1	vehicle	kind	seats	wheels?
2		car	5	<input checked="" type="checkbox"/>
3		trolley	40	<input checked="" type="checkbox"/>
4		trolley		<input checked="" type="checkbox"/>
5		truck	3	<input checked="" type="checkbox"/>
6		helicopter	6	<input type="checkbox"/>
7		bus	42	<input checked="" type="checkbox"/>
8		tractor	1	<input checked="" type="checkbox"/>
9		boat	4000	<input type="checkbox"/>
10		scooter	2	<input checked="" type="checkbox"/>
11		bike	1	<input checked="" type="checkbox"/>

[Access this spreadsheet on Google Sheets](#) ➤

- **NB.:** On export, **data validation rules** (e.g., number ranges, value constraints ☒☐) , formulæ, or cell links **are lost**.

4 : Importing Data Into Tables, Exporting Query Results Into Files

1. Save repeated imports, use **COPY** to **copy a file into a table**:

```
COPY table  
FROM path_to_file (configuration parameters)
```

- Can now query *table* like a regular database table.

2. **Copy SQL query results into a file** external to the DBMS:


```
COPY (query)  
TO path_to_file (configuration parameters)
```

The *configuration parameters* determine file format (e.g., CSV), delimiters, header output, compression, overwrite/append mode, ...

Using DuckDB as a CSV Processor

DBMS may act as a processing engine for database-external data:

- + SQL is versatile, expressive, and executed efficiently.
- Need to pay for data import and export.



 #010

```

❶  D COPY (SELECT ...
      FROM   input_file
      WHERE  ...
      ORDER BY ...)
    TO output_file;
  
```

```

❷  D COPY table
      FROM input_file;
    D COPY (SELECT ...
      FROM   table
      WHERE  ...
      ORDER BY ...)
    TO output_file;
  
```



5 | Is CSV Too Simple? Ubiquitous, But Fragile

The CSV format is documented in [RFC 4180](#) . Still, CSV files in daily practice deviate in a variety of ways¹, including

- choice of **delimiters** (frequent: `,` `;` `|` (pipe) `␣` (tab)),
- **quoting** (enclose columns with significant spaces `_` in `"` or `'`),
- character **escape** sequences (what if `,` `"` `␣` are in columns?),
- absence/presence of **headers or preamble text** above data rows,
- choice of representation of **missing values** (e.g., `...`, `,,`, `...`), or
- a **mixture of data types** in a column.

DuckDBs aims to auto-adapt to these CSV dialects using a “multi-hypothesis sniffing” algorithm².

¹ DuckDB's CSV reader comes with 25+ configuration options. To get an impression of the variety of CSV dialects out there, see the paper [Characteristics of Open Data CSV Files](#)  (2016).

² See the blog post [DuckDB's CSV Sniffer: Automatic Detection of Types and Dialects](#)  and the paper [Multi-Hypothesis CSV Parsing](#)  (2017) by Hannes Mühleisen *et al.*

DuckDB's CSV “Sniffer”

1. **Sample 20480 lines** from all over the file (unless we need to read top to bottom, e.g., if input is `stdin` or compressed).
2. Try **CSV dialects** based on 24 selected combinations of the options below. Choose the CSV dialect options that lead to a *consistent and maximum number of columns* per row:

parameter	options
<code>delim</code>	<code>,</code> <code> </code> <code>␣</code> <code>␣</code> <code>;</code>
<code>quote</code>	<code>"</code> <code>'</code> (empty)
<code>escape</code>	<code>"</code> <code>'</code> <code>\</code> (empty)

3. **Detect types.** Try to cast values to candidate types *in order*: `[NULL,)` `boolean`, `int`, `double`, `time`, `date`, `timestamp`, `text`].
4. **Header detection:** Do columns in first row match detected types (*is data*, thus generate own header) or not (*is header*)?


DuckDB's CSV “Sniffer”: Dialect and Type Detection

Dialect detection:

 #011

 flights.csv	# columns/delimiter			
	,		^H _T	;
FlightDate Carrier Origin Destination ^N _L	1	4	1	1
1988-01-01 AA New York, NY Los Angeles, CA ^N _L	3	4	1	1
1988-01-02 AA New York, NY Los Angeles, CA ^N _L	3	4	1	1
1988-01-03 AA New York, NY Los Angeles, CA ^N _L	3	4	1	1
		↑		

Type detection (file reading order →):

 band.csv	candidate types for	
	column 0	column 1
Name, Age	(skip)	(skip)
,	[NULL,boolean,...,text]	[NULL,boolean,...,text]
Mike Lindup, 65	[text]	[int,double,...,text]
Mark King, 66.2	[text]	[double,...,text]
	▼ ↑	↑